Running Deep Learning Experiments

Last modified: 2026-11-04

```
# The Python chunks on this page import these; reticulate provisions them.
# Attaching reticulate gives the inline R its `py` object.
library(reticulate)
source("../tools/use-torch.R")
use_torch(c("numpy", "pandas", "matplotlib"))
```

Brian Hutchinson's CSCI 481/581 sets the order of topics in these notes, which follow it with his permission; this page covers his module on running deep learning experiments. The exercises are new.

1 Welcome

1.1 Plan for this lecture

Earlier chapters trained networks with settings picked for them: a learning rate[1], a width[2], a number of epochs[3]. This lecture is about choosing such settings with experiments, and trusting what the experiments say:

1. how much a score changes when only the random seed changes;
2. how to compare two widths[4] fairly when the best learning rate[5] depends on the width[6];
3. how to lay out a sweep of learning rates[7], and check that it was wide enough;
4. when training longer would help;
5. what it takes to run the same experiment twice and get the same numbers.

lds's chapter on machine learning experiments[8] covers the parts that hold for any model: the training, validation and test sets[9], data snooping, and grid and random search. This lecture adds what is particular to training deep networks.

1.2 The digits

```
import pandas as pd
import torch
from torch import nn
import matplotlib.pyplot as plt

digits = pd.read_csv("data/optdigits-test.csv", header=None)
x = torch.tensor(digits.iloc[:, :64].to_numpy(), dtype=torch.float32) / 16
y = torch.tensor(digits[64].to_numpy())
order = torch.randperm(len(x), generator=torch.Generator().manual_seed(481))
```

[1] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[2] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[3] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[4] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[5] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[6] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[7] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[8] https://morrison-lab.github.io/lds/chapters/ml-experimentation.html
[9] https://morrison-lab.github.io/lds/chapters/big-ideas.html#def-test-set

```
train_idx, val_idx = order[:500], order[500:1000]
n_train, n_val = len(train_idx), len(val_idx)
n_unused = len(x) - n_train - n_val
loss_fn = nn.CrossEntropyLoss()


def train_and_score(width, lr, seed, n_epochs=20, batch_size=50):
    # One training run: an MLP with one hidden layer of `width` ReLU units,
    # trained by minibatch SGD. Returns the validation accuracy after each epoch.
    torch.manual_seed(seed)
    net = nn.Sequential(nn.Linear(64, width), nn.ReLU(), nn.Linear(width, 10))
    optimizer = torch.optim.SGD(net.parameters(), lr=lr)
    shuffler = torch.Generator().manual_seed(seed)
    scores = []
    for epoch in range(n_epochs):
        for batch in torch.randperm(len(train_idx), generator=shuffler).split(batch_size):
            rows = train_idx[batch]
            optimizer.zero_grad()
            loss_fn(net(x[rows]), y[rows]).backward()
            optimizer.step()
        with torch.no_grad():
            correct = net(x[val_idx]).argmax(dim=1) == y[val_idx]
        scores.append(correct.float().mean().item())
    return scores


print(n_train, n_val, n_unused)
```

```
500 500 797
```

The exercises use the 1,797 digits of `data/optdigits-test.csv`[10] [1], each an 8 by 8 image with 64 pixels. 500 random images form the training set[11] and another 500 the validation set[12]. The other 797 stay unused, as a test set[13] would until the model is final. `train_and_score()` builds an MLP with one hidden layer of `width` ReLU units. It trains the MLP by minibatch[14] SGD[15] with learning rate[16] `lr`, 50 images to a minibatch, for 20 epochs[17] unless asked for more. It returns the accuracy[18] on the validation set after each epoch[19].

2 Settings and seeds

2.1 What training does not learn

Exercise 2.1 (Settings that training does not learn)**.** List the numbers that `train_and_score()` uses but SGD[a] never changes.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-sgd

[10] https://morrison-lab.github.io/dds/chapters/data/optdigits-test.csv
[11] https://morrison-lab.github.io/lds/chapters/big-ideas.html#def-training-set
[12] https://morrison-lab.github.io/lds/chapters/big-ideas.html#def-validation-set
[13] https://morrison-lab.github.io/lds/chapters/big-ideas.html#def-test-set
[14] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[15] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-sgd
[16] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[17] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[18] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[19] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch

Solution

Solution 2.1. SGD[a] changes only the weights and biases of the two `nn.Linear` layers. These numbers are set before training starts, and stay as set:

- the width[b] of the hidden layer, `width`;
- the learning rate[c], `lr`;
- the number of epochs[d], `n_epochs`;
- the minibatch[e] size, `batch_size`;
- the random seed, `seed`;
- the sizes of the training and validation sets[f], 500 and 500.

The first four are hyperparameters[g]: settings that we choose, rather than learn from the data. So are choices that are not numbers, such as SGD itself, the ReLU activation and the single hidden layer. The seed is not a setting of that kind. No seed suits new data better than another, so a seed chosen for its validation score has been fitted to noise. The set sizes are a choice about the data, not the model.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-sgd
[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[c] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[d] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[e] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[f] https://morrison-lab.github.io/lds/chapters/big-ideas.html#def-validation-set
[g] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

2.2 The same settings, ten seeds

Exercise 2.2 (Ten seeds)**.** Train a network of width[a] 64 with learning rate[b] 0.1 once for each seed from 0 to 9. How much does the final validation accuracy[c] vary?

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

Solution

Solution 2.2.

```
wide_scores = torch.tensor([train_and_score(64, 0.1, seed)[-1] for seed in range(10)])
seed_low, seed_high = wide_scores.min().item(), wide_scores.max().item()
seed_sd = wide_scores.std().item()
print(wide_scores)
```

```
tensor([0.8900, 0.8860, 0.8740, 0.8960, 0.8760, 0.8920, 0.9040, 0.8900, 0.8900,
        0.8960])
```

```
print(seed_low, seed_high, seed_sd)
```

```
0.8740000128746033 0.9039999842643738 0.00904556643217802
```

The ten validation accuracies[a] run from 0.874 to 0.904, with standard deviation 0.009. Only the seed changed between the runs. The seed sets the starting weights and the order of the minibatches[b]. So two runs can differ by about 0.03 in validation accuracy[c] with every hyperparameter[d] the same.

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[d] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

2.3 One run of each

Exercise 2.3 (One run of each)**.**

```
narrow_seed0 = train_and_score(32, 0.1, 0)[-1]
wide_seed0 = wide_scores[0].item()
print(narrow_seed0, wide_seed0)
```

```
0.8820000290870667 0.8899999856948853
```

With seed 0 and learning rate[a] 0.1, a network of width[b] 32 reaches validation accuracy[c] 0.882, and one of width 64 reaches 0.890. Does this show that width 64 is better at this learning rate?

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

Solution

Solution 2.3.

```
narrow_scores = torch.tensor([train_and_score(32, 0.1, seed)[-1] for seed in range(10)])
mean_gap = (wide_scores.mean() - narrow_scores.mean()).item()
gap_se = ((wide_scores.var() + narrow_scores.var()) / 10).sqrt().item()
print(narrow_scores.mean().item(), wide_scores.mean().item(), mean_gap, gap_se)
```

```
0.8855999708175659 0.8894000053405762 0.003800034523010254 0.004646146669983864
```

```
# Stop the render if the prose below stops being true.
assert abs(wide_seed0 - narrow_seed0) < seed_high - seed_low
assert abs(mean_gap) < 2 * gap_se
```

No. The two runs differ by 0.008, less than the 0.030 that the seed alone moved width[a] 64 in Solution 2.2. With ten seeds of each, the mean accuracy[b] of width 64 is higher by 0.0038. That difference has standard error[c] 0.0046, so it is less than two standard errors from 0, and ten seeds each do not settle the question either. A comparison of single runs measures the two seeds as much as the two widths[d].

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[c] https://morrison-lab.github.io/lds/chapters/linear-regression.html#def-standard-error
[d] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width

3 Comparing settings fairly

3.1 Two widths

Exercise 3.1 (A fair comparison)**.** A classmate wants to know whether a network of width[a] 512 classifies the digits better than one of width 8. They train one of each with learning rate[b] 0.1. What could make that comparison unfair?

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

Solution

Solution 3.1. Nothing says that one learning rate[a] is best for both widths[b]. If 0.1 suits width[c] 512 and is too small for width 8, the comparison measures how well each width suits 0.1, as well as the width itself. A fair comparison tunes the learning rate for each width on its own, and compares each width at its own best rate. One run of each has the problem of Solution 2.3 too, so each setting gets several seeds.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[c] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width

3.2 A sweep for each width

Exercise 3.2 (A learning-rate sweep)**.** For each width[a], 8 and 512, train with each learning rate[b] in 0.003, 0.01, 0.03, 0.1, 0.3, 1 and 3, and with seeds 0, 1 and 2. How much higher is width 512's mean validation accuracy[c] than width 8's, when each width uses the learning rate that suits it best?

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

Solution

Solution 3.2.

```
rates = [0.003, 0.01, 0.03, 0.1, 0.3, 1, 3]
sweep = pd.DataFrame(
    [
        {"width": width, "lr": lr, "seed": seed, "accuracy": train_and_score(width, lr, seed)[-1]}
        for width in [8, 512]
        for lr in rates
        for seed in range(3)
    ]
)
mean_accuracy = sweep.groupby(["width", "lr"])["accuracy"].mean().unstack("width")
best_rate = mean_accuracy.idxmax()
best_accuracy = mean_accuracy.max()
best_rate_8, best_rate_512 = best_rate[8], best_rate[512]
edge_accuracy_512 = mean_accuracy.loc[3, 512]
gap_at_01 = mean_accuracy.loc[0.1, 512] - mean_accuracy.loc[0.1, 8]
gap_at_best = best_accuracy[512] - best_accuracy[8]
print(mean_accuracy.round(3))
```

```
width    8      512
lr
0.003  0.166  0.410
0.010  0.237  0.679
0.030  0.299  0.875
0.100  0.719  0.903
0.300  0.879  0.929
1.000  0.808  0.949
3.000  0.091  0.088
```

```
print(best_rate.to_dict(), gap_at_01, gap_at_best)
```

```
{8: 0.3, 512: 1.0} 0.18400001525878906 0.07066665093104052
```

```
# Stop the render if the prose below stops being true.
assert best_rate[8] < best_rate[512] < max(rates)
assert 2 * seed_sd < gap_at_best < gap_at_01
# Stop the render if the figure's alt text stops being true.
assert best_rate[8] == 0.3 and best_rate[512] == 1
assert abs(best_accuracy[512] - 0.95) < 0.015
assert abs(best_accuracy[8] - 0.88) < 0.015
assert mean_accuracy.loc[3].max() < 0.15
```

```
fig, ax = plt.subplots(figsize=(6, 3.5))
for width, marker in [(8, "o"), (512, "s")]:
    runs = sweep[sweep["width"] == width]
    ax.scatter(runs["lr"], runs["accuracy"], marker=marker, alpha=0.5, label=f"width {width}");
    ax.plot(mean_accuracy.index, mean_accuracy[width]);
ax.set_xscale("log");
ax.set_xlabel("learning rate");
ax.set_ylabel("validation accuracy");
ax.legend();
fig.tight_layout();
plt.show()
```

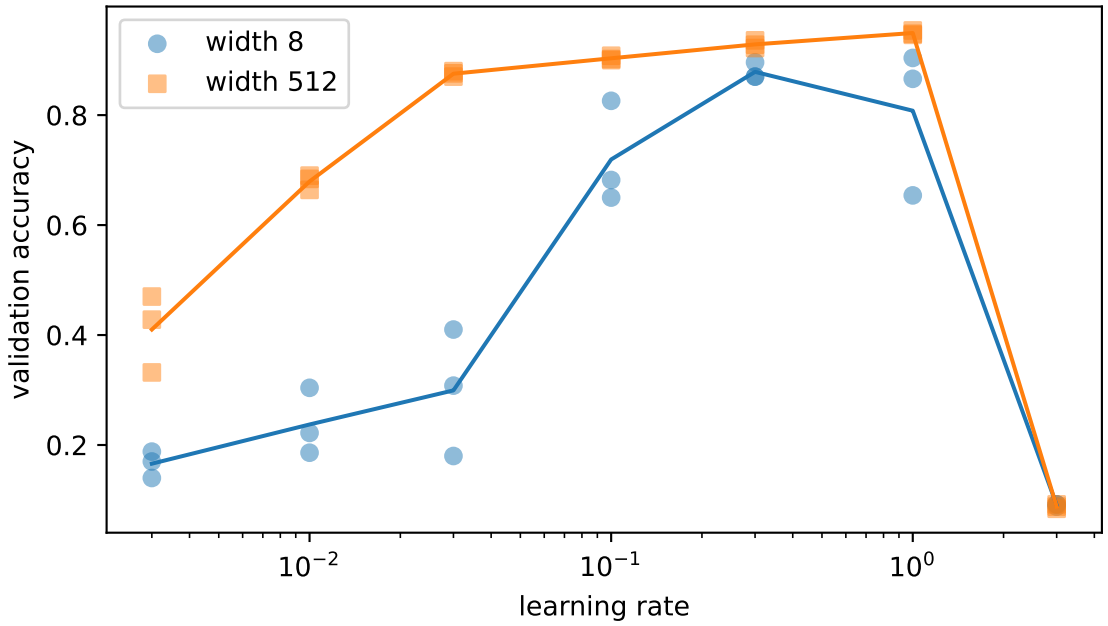


Figure 1: Final validation accuracy[a] of each run in the sweep, and the mean over the three seeds at each learning rate[b].

Width[c] 8 does best with learning rate[d] 0.3, and width 512 with 1. At learning rate 0.1, width 512's mean accuracy[e] is higher by 0.184. With each width at its own best rate, width 512's mean accuracy is higher by 0.071, 8 times the seed-to-seed standard deviation of Solution 2.2. So the wider network does classify better, by 38% of what the comparison at one rate said.

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[d] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[e] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

3.3 Three roles

Definition 3.1 (Scientific hyperparameter)**.** A **scientific hyperparameter** of an experiment is a hyperparameter[a] whose effect on the model's performance the experiment is meant to measure [2].

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

Example 3.1 (The width in the sweep)**.** In Solution 3.2, the width[a] is the scientific hyperparameter: the sweep asks whether width 512 classifies the digits better than width 8.

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width

Definition 3.2 (Nuisance hyperparameter)**.** A **nuisance hyperparameter** of an experiment is a hyperparameter[a] that the experiment tunes for each value of its scientific hyperparameters, so that those values are compared fairly [2].

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

Example 3.2 (The learning rate in the sweep)**.** In Solution 3.2, the learning rate[a] is a nuisance hyperparameter. The sweep tunes it for each width[b], 8 and 512, and compares the widths[c] at their own best rates.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[c] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width

Definition 3.3 (Fixed hyperparameter)**.** A **fixed hyperparameter** of an experiment is a hyperparameter[a] that keeps one value throughout the experiment [2]. The experiment's conclusions hold for that value, and may not hold for others.

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

Example 3.3 (What the sweep held fixed)**.** In Solution 3.2, the fixed hyperparameters include the minibatch[a] size of 50, the 20 epochs[b], the single hidden layer and the choice of SGD[c]. So the sweep shows that width[d] 512 beats width 8 after 20 epochs of SGD with minibatches[e] of 50. It does not show it for 5 epochs, or for Adam[f].

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[c] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-sgd
[d] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[e] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[f] https://morrison-lab.github.io/mds/optimization-gradient-descent.html#def-adam

Remark (A role belongs to an experiment)

Remark 3.1 (A role belongs to an experiment). One hyperparameter[a] can play a different role in each experiment [2]. The learning rate[b] is a nuisance hyperparameter in Solution 3.2. An experiment that asks whether learning-rate warmup helps makes the warmup's length a scientific hyperparameter, and still tunes the learning rate as a nuisance. An experiment that cannot afford to tune the learning rate makes it a fixed hyperparameter, and its conclusions then come with that caveat [2].

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

4 Laying out a sweep

4.1 Spacing the rates

Exercise 4.1 (Spacing the rates)**.** Each learning rate[a] in Exercise 3.2 is about 3 times the one before it. What would go wrong with seven rates spread evenly from 0.003 to 3?

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

Solution

Solution 4.1.

```
import numpy as np

even_rates = np.linspace(0.003, 3, 7)
print(even_rates.round(4))
```

```
[0.003  0.5025 1.002  1.5015 2.001  2.5005 3.    ]
```

Evenly spread rates go up by 0.4995 at each step, so the second is already 0.5025. Nothing between 0.003 and 0.5025 would be tried, and that interval holds width[a] 8's best learning rate[b], 0.3, and every rate at which width 512 was still improving in Figure 1. Good learning rates[c] can differ by factors of 10 or more, so a sweep spaces them by a constant factor: evenly on a log scale. Random search draws a learning rate the same way, uniform on its logarithm (lds[d]).

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[d] https://morrison-lab.github.io/lds/chapters/ml-experimentation.html#tuning

4.2 The edge of the sweep

Exercise 4.2 (The edge of the sweep)**.** Suppose the sweep in Solution 3.2 had stopped at learning rate[a] 1. What would its result for width[b] 512 leave unknown?

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width

Solution

Solution 4.2. Width[a] 512's best learning rate[b] would then be the largest one tried. A larger rate might do better still, and the sweep could not say. When the best value sits at the edge of the range searched, the range is suspect, and the next step is to extend it in that direction [2]. The actual sweep went one step further, to 3, where width 512's mean validation accuracy[c] drops to 0.088, about the 0.1 that guessing a digit at random would score. So the rate that suits width 512 best lies between 0.3 and 3, and a finer sweep in that range could pin it down.

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

5 How long to train

5.1 More epochs

Exercise 5.1 (More epochs)**.** Would width[a] 512, at the learning rate[b] that suited it best in Solution 3.2, reach a higher validation accuracy[c] with 100 epochs[d] of training instead of 20?

[a] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[c] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[d] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch

Solution

Solution 5.1.

```
long_run = train_and_score(512, best_rate[512], 0, n_epochs=100)
at_20, best_long = long_run[19], max(long_run)
best_epoch = long_run.index(best_long) + 1
print(at_20, best_long, best_epoch, long_run[-1])
```

```
0.9480000138282776 0.9580000042915344 31 0.9539999961853027
```

```
# Stop the render if the prose below stops being true.
assert best_long - at_20 < seed_high - seed_low
assert min(long_run[19:]) > at_20 - 0.01
# Stop the render if the figure's alt text stops being true.
assert abs(at_20 - 0.95) < 0.015
assert all(abs(score - at_20) < 0.015 for score in long_run[19:])
```

```
fig, ax = plt.subplots(figsize=(6, 3))
ax.plot(range(1, 101), long_run);
ax.axvline(20, linestyle="--", color="gray");
ax.set_xlabel("epoch");
ax.set_ylabel("validation accuracy");
fig.tight_layout();
plt.show()
```

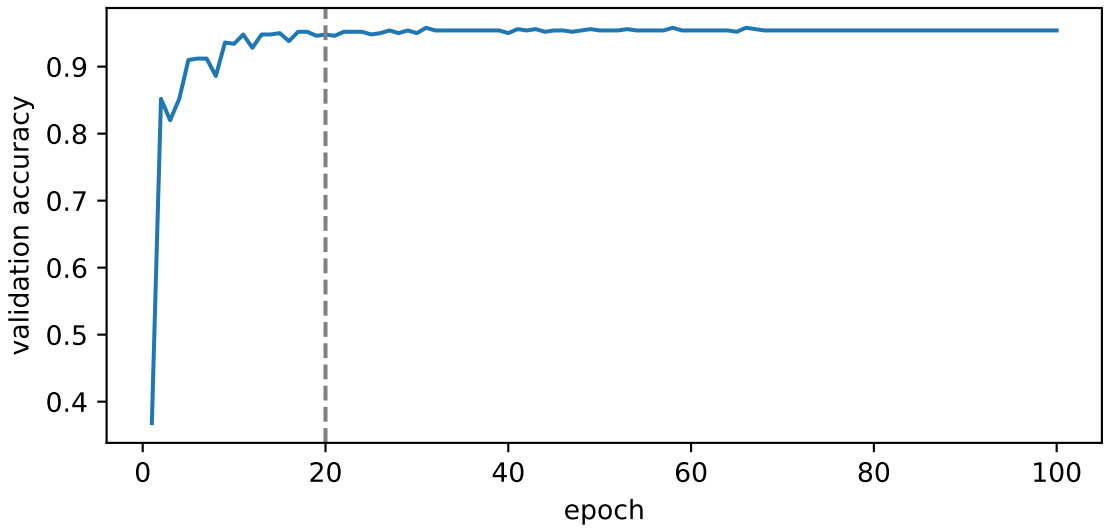


Figure 2: Validation accuracy[a] after each epoch[b] of one 100-epoch run of width[c] 512 at learning rate[d] 1; the dashed line marks epoch 20.

Hardly. After 20 epochs[e] the validation accuracy[f] is 0.948. Its highest value over 100 epochs is 0.958, at epoch[g] 31, less above the 20-epoch value than the seed alone moved width[h] 64 in Solution 2.2. With the learning rate[i] the sweep tuned for 20 epochs, training 5 times as long buys nothing that this experiment can detect.

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[c] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[d] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[e] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[f] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[g] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[h] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[i] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

Definition 5.1 (Compute-bound)**.** Training is **compute-bound** when how long we are willing to train limits the model's performance, rather than the amount of data or anything else [2]. Then training longer, with the hyperparameters[a] tuned for the longer run, improves the validation performance.

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters

Example 5.1 (A slow rate is not the problem)**.**

```
slow_run = train_and_score(512, 0.03, 0, n_epochs=100)
long_rates = [0.1, 0.3, 1]
long_scores = [train_and_score(512, lr, 0, n_epochs=100)[-1] for lr in long_rates]
best_long_tuned = max(long_scores)
print(slow_run[19], slow_run[-1], long_scores)
```

```
0.8799999952316284 0.9139999747276306 [0.9459999799728394, 0.9440000057220459, 0.9539999961853027]
```

```
# Stop the render if the prose below stops being true.
assert slow_run[-1] - slow_run[19] > seed_high - seed_low
assert slow_run[-1] < at_20
assert all(s < l for s, l in zip(slow_run[19:], long_run[19:]))
assert best_long_tuned - at_20 < seed_high - seed_low
# Stop the render if the figure's alt text stops being true.
assert abs(slow_run[19] - 0.88) < 0.015
assert abs(slow_run[-1] - 0.91) < 0.015
```

```
fig, ax = plt.subplots(figsize=(6, 3))
ax.plot(range(1, 101), long_run, label=f"learning rate {best_rate[512]}");
ax.plot(range(1, 101), slow_run, label="learning rate 0.03");
ax.axvline(20, linestyle="--", color="gray");
ax.set_xlabel("epoch");
ax.set_ylabel("validation accuracy");
ax.legend();
fig.tight_layout();
plt.show()
```

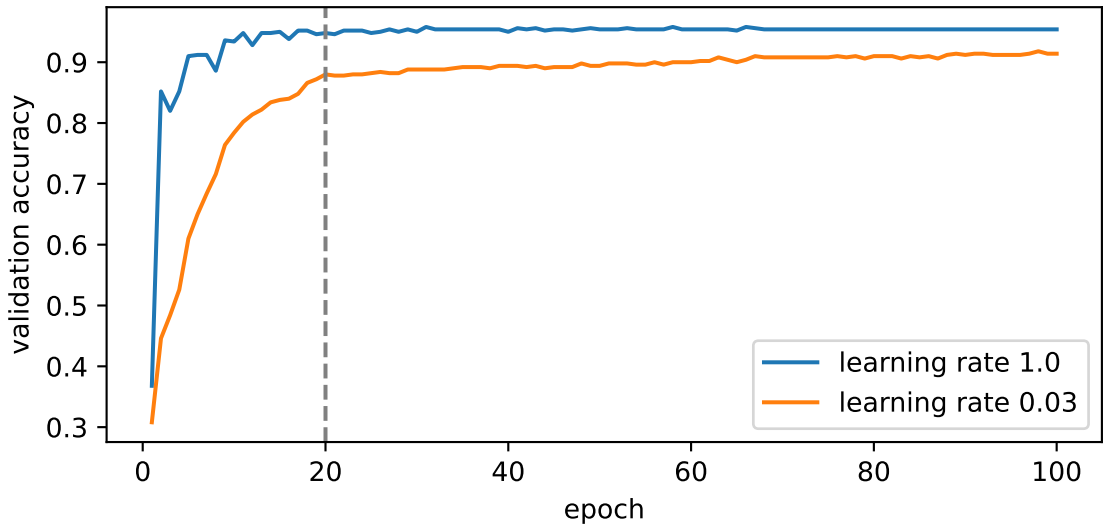


Figure 3: Validation accuracy[a] after each epoch[b] of 100-epoch runs of width[c] 512 at two learning rates[d]; the dashed line marks epoch 20.

At learning rate[e] 0.03, width[f] 512's validation accuracy[g] rises from 0.880 after 20 epochs[h] to 0.914 after 100. That run alone would suggest that 20 epochs hold the model back. But at its best learning rate the network gets 0.948 in 20 epochs, more than the slow run gets in 100. Tuned for 100 epochs instead, over learning rates[i] 0.1, 0.3 and 1, the best 100-epoch run reaches 0.954, less above the 20-epoch run than the seed alone moved width 64 in Solution 2.2. Longer training does not help, so training this network on these 500 images is not compute-bound. Something other than training time limits it.

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[c] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[d] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[e] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate
[f] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[g] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy
[h] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[i] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

Remark (Spending a fixed budget)

Remark 5.1 (Spending a fixed budget). The sweep in Solution 3.2 trained 42 networks for 20 epochs[a] each, 840 epochs in all. One 100-epoch run of one network takes as many epochs as 5 of those trials. Here, where training is not compute-bound, the extra 80 epochs in Solution 5.1 gained less than a change of seed moves the score, and the same 80 epochs would have paid for 4 more trials of the same width[b], such as more seeds or more learning rates[c]. When training is compute-bound, Godbole et al. [2] still advise against making every trial as long as the budget allows. They tune in rounds, starting with many short runs to find good settings and ending with a few long runs at the best of them.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch
[b] https://morrison-lab.github.io/lds/chapters/neural-nets.html#def-width
[c] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-learning-rate

6 Running it again

6.1 Where the randomness comes from

Exercise 6.1 (Where the randomness comes from)**.** List every random draw that a run of `train_and_score()` depends on.

Solution

Solution 6.1.

- **The starting weights.** `torch.manual_seed(seed)` seeds PyTorch's global random number generator, and each `nn.Linear` draws its starting weights and biases from it.
- **The order of the minibatches**[a]**.** `shuffler`, a generator of its own seeded with `seed`, draws a new order of the training images in each epoch[b].
- **The split.** `order`, drawn once outside the function from a generator seeded with 481, decides which images are for training and which for validation. Every run shares it, so the runs differ in their seeds and settings but not in their data.

A network with dropout would draw one more random mask at each step, from the global generator.

[a] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-minibatch
[b] https://morrison-lab.github.io/lds/chapters/gradient-descent.html#def-epoch

6.2 The same run twice

Exercise 6.2 (The same run twice)**.** Does calling `train_and_score(64, 0.1, 0)` twice give the same validation accuracies[a]?

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

Solution

Solution 6.2.

```
first, second = train_and_score(64, 0.1, 0), train_and_score(64, 0.1, 0)
print(first == second, first[-1])
```

```
True 0.8899999856948853
```

```
print(torch.__version__)
```

```
2.15.0+cu132
```

```
assert first == second
```

Yes: all 20 validation accuracies[a] agree, because each call reseeds both generators of Solution 6.1 before it draws anything. That holds on one machine, with one version of each library. The printed `torch.__version__` is the PyTorch version that rendered this page.

[a] https://morrison-lab.github.io/lds/chapters/metrics.html#def-accuracy

Remark (Seeds are not enough)

Remark 6.1 (Seeds are not enough). PyTorch's documentation warns that completely reproducible results are not guaranteed across PyTorch releases or platforms, or between a CPU and a GPU, even with identical seeds [3]. Some GPU operations are nondeterministic, and `torch.use_deterministic_algorithms(True)` makes PyTorch choose deterministic ones, often more slowly, or raise an error where it has none [3]. So record more than the seed: the hyperparameters[a], the versions of the code, the data and the libraries, and the trained weights themselves. lds's section on tracking runs[b] gives the habits that make a result traceable a month later.

[a] https://morrison-lab.github.io/lds/chapters/model-selection.html#sec-hyperparameters
[b] https://morrison-lab.github.io/lds/chapters/ml-experimentation.html#tracking-runs

6.3 Further reading

Read

The Deep Learning Tuning Playbook[a] [2] is the source of this lecture's vocabulary, and covers choosing a batch size and a starting configuration as well. *Deep Learning*'s chapter 11[b] covers manual and automatic hyperparameter search [4, Sec. 11.4].

[a] https://github.com/google-research/tuning_playbook
[b] https://www.deeplearningbook.org/contents/guidelines.html

7 References

[1] E. Alpaydin and C. Kaynak, "Optical recognition of handwritten digits." UCI Machine Learning Repository, 1998. doi: 10.24432/C50P49[20].

[2] V. Godbole, G. E. Dahl, J. Gilmer, C. J. Shallue, and Z. Nado, "Deep learning tuning playbook." GitHub repository, 2023. Available: https://github.com/google-research/tuning_playbook

[3] PyTorch Contributors, "Reproducibility." PyTorch documentation, 2026. Available: https://docs.pytorch.org/docs/stable/notes/randomness.html

[4] I. Goodfellow, Y. Bengio, and A. Courville, *Deep learning*. MIT Press, 2016. Available: https://www.deeplearningbook.org/

[20] https://doi.org/10.24432/C50P49